

Setting up the Environment -

Airflow Download and Setup Steps:

-Download Virtual Box and install

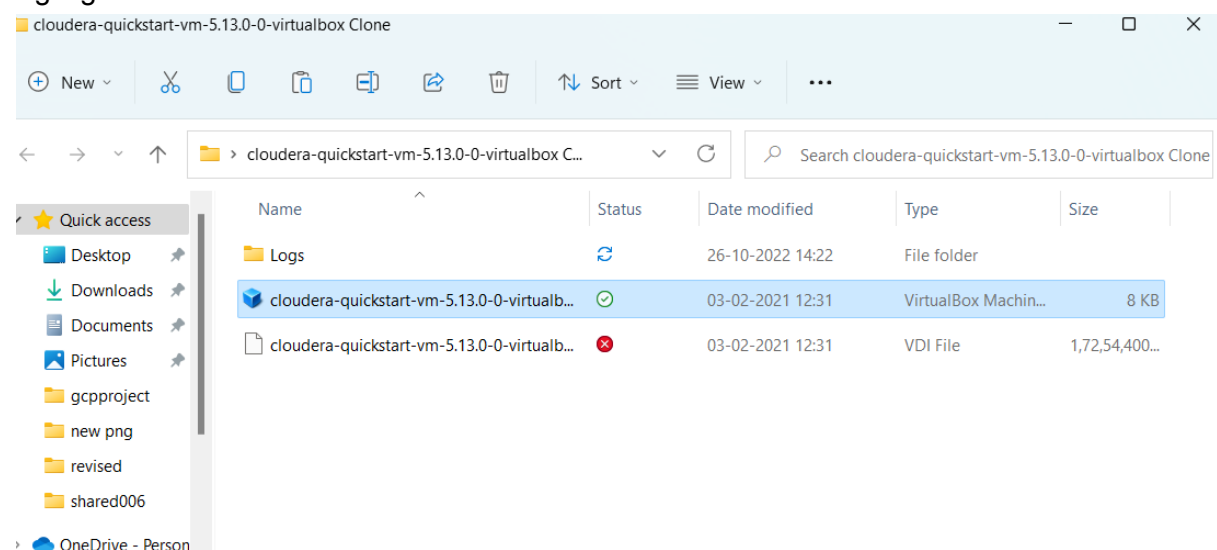
<https://www.virtualbox.org/wiki/Downloads>

Select the respective host Operating System and download the file and install.


-We are providing a VM which has airflow installed called a clone VM. **Just download , import the VM and start using it as steps given below**

<https://drive.google.com/drive/folders/1fXViofZOqdvX2bKY3RVwyTpCbgvFS7il?usp=sharing>

-Extract the files from the zip. Open the cloudera-quickstart-vm folder that gets created after extracting the zip. Double click on the cloudera-quickstart-vm-5.13.0-0-virtualb... as highlighted in the screenshot below.



This will open the file in the virtual box. Select the important VM and click on **Start** as shown below.

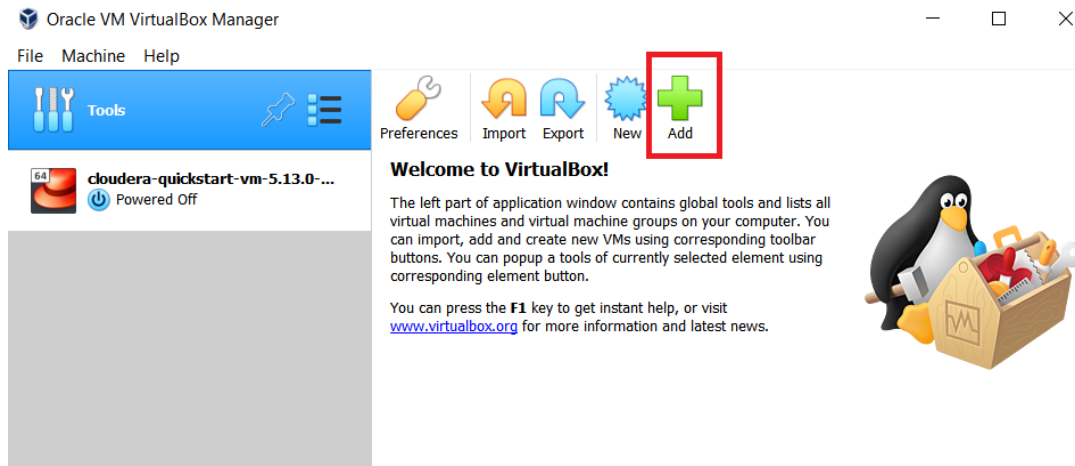


[**ERRORS & SOLUTIONS**

If you are unable to import the VM and if you see a .vdi instead of .ovf after extracting, follow the steps below to import the VM to Virtual box

Either right click on the .vdi file and open with Virtual box

OR



Click on the **+ADD** option in the Virtual Box and select the .vdi file. This should start the import with default values.

]

Additional Background Tasks:

1. Create a **shared folder** as explained in the course to transfer the files from host to cloudera VM. You could refer to the 3 General Things section of Week1 for shared folder creation steps.
2. Drive Link to Download BigData_Project_Demo Folder. This contains all the scripts used in the project.

<https://drive.google.com/drive/folders/1eWgaOZN54aePfTbtIm8oeawjFkEY-s?usp=sharing>

[Note: If you downloaded the BigData_Project_Demo on your host windows or mac machine, transfer it to cloudera VM using Shared folder]

Follow below steps to setup airflow

1. **su** (pwd: cloudera)
2. **cd airflow**
3. **source venv/bin/activate**
4. **export AIRFLOW_HOME=\$PWD**
5. Create dag folder if not exists
[mkdir dags]
6. Create python file inside dag folder(for the python file, copy paste the export_card_txns.py from BigData_Project_Demo folder, present under Airflow_Scripts)

7. Run the command **"airflow webserver -p 8080"**
8. Open another terminal — do steps 1-4 — and Run the command **"airflow scheduler"**
9. Go to Browser — localhost:8080 — you can see your dag with dag id not with python file name

[**ERRORS & SOLUTIONS:**

1.If you run into errors while starting the airflow server, run the following on the same terminal -

```
pip install wheel
pip install --upgrade pip
pip install sshtunnel
```

2.If you run into PARAMIKO error, execute below commands:

```
export CRYPTOGRAPHY_DONT_BUILD_RUST=1
pip3 install --upgrade homeassistant
pip install paramiko (or pip3 install paramiko)
```

After executing this, if you encounter error of ssh tunnel, run below command:

```
pip install sshtunnel
```

3.If you encounter error: Already running on PID , execute below command:

```
sudo lsof -i tcp:8080 (so you will get list of running pid's, now execute below
kill command to kill all the mentioned pid's)
```

```
kill -9 <pid>
```

Also remove the airflow-webserver.pid by executing below command:

```
rm airflow-webserver.pid
```

]

Tasks:

The following Datasets required for the project are present in the Project Datasets folder in the above google drive link.

1. Card Transactions History Data - card_transactions.csv
2. Member Score Dataset
3. Member Details Dataset

Task1: Copy “card_transactions.csv” file from local system to HDFS

```
hadoop fs -mkdir project_input_data
hadoop fs -put Desktop/card_transactions.csv project_input_data/
hadoop fs -cat project_input_data/card_transactions.csv | wc -l
```

Connect to mysql
mysql -u root -p
Password :cloudera

Task2: Creating tables in MySQL (Run in MySQL terminal)

```
create table stg_card_transactions (
card_id bigint,
member_id bigint,
amount int,
postcode int,
pos_id bigint,
transaction_dt varchar(255),
status varchar(50)
);
```

```
create table card_transactions (
card_id bigint,
member_id bigint,
amount int,
postcode int,
pos_id bigint,
transaction_dt datetime,
status varchar(50),
PRIMARY KEY(card_id, transaction_dt)
);
```

Task3: Sqoop export to the card_transactions table in MySQL database for card_transactions.csv(Using Airflow) and delete the file from HDFS.

Navigate to -> BigData_Project_Demo -> Airflow_Scripts
copy : export_card_txns.py

Navigate to Home Directory -> airflow -> dags
paste : export_card_txns.py

Encrypting MySQL Password - (Run in the terminal)

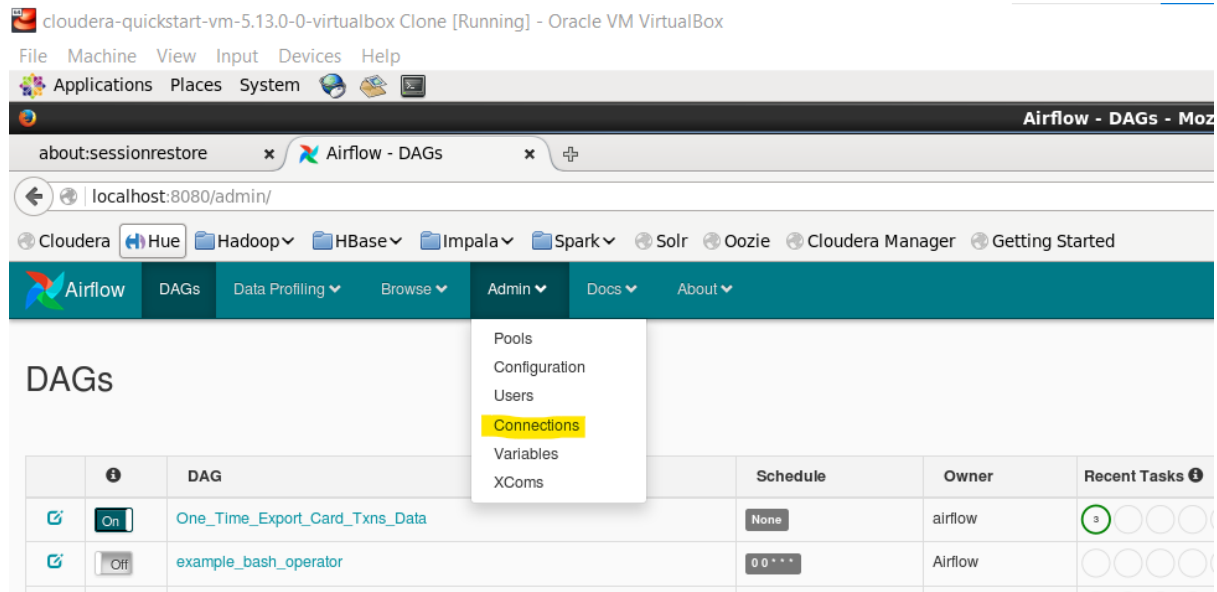
```
hadoop credential create mysql.bigdatapassword -provider
jceks://hdfs/user/cloudera/mysql.dbpassword.jceks
```

Copy the sh file (sqoop_export_card_txns.sh file to Desktop and give permission to this file :
cd Desktop
chmod 777 sqoop_export_card_txns.sh)

Creating Variables and Connections

After having started the Airflow server, When you traverse to localhost:8080 , the DAGS will be visible if the server started successfully

Under the **Admin** tab ->



Setting up Connection

click on **Connections**:

Under **create** tab, fill in the following details and save the connection

Conn id : cloudera

Conn Type : SSH

Host : quickstart.cloudera

Username : cloudera

Password: cloudera

localhost:8080/admin/connection/edit/?id=39&url=%2Fadmin%2Fconnection%2F

Cloudera Hue Hadoop HBase Impala Spark Solr Oozie Cloudera Manager Getting Started

Airflow DAGs Data Profiling Browse Admin Docs About

Connection [edit]

List Create Edit

Conn Id * cloudera

Conn Type SSH

Host quickstart.cloudera

Username cloudera

Password

Port

Extra

Save Save and Add Another Save and Continue Editing Cancel

And Save

Setting up Variable

click on **Variables**:

Under **create** tab, fill in the following details and save the connection

Key : card_txns_export_shell_command

Val : `./sqoop_export_card_txns.sh quickstart.cloudera:3306 bigdatapipeline root stg_card_transactions`

localhost:8080/admin/variable/edit/?id=1&url=%2Fadmin%2Fvariable%2F

Cloudera Hue Hadoop HBase Impala Spark Solr Oozie Cloudera Manager Getting Started

Airflow DAGs Data Profiling Browse Admin Docs About

Variable [edit]

List Create Edit

Key * card_txns_export_shell_command

Val ./sqoop_export_card_txns.sh quickstart.cloudera:3306 bigdatapipeline root stg_card_transactions

Save Save and Add Another Save and Continue Editing Cancel

And Save

Select **One_Time_Export_Card_Txns_Data** dag

localhost:8080/admin/

Search

Cloudiera

Hue

Hadoop

HBase

Impala

Spark

Solr

Oozie

Cloudiera Manager

Getting Started

Airflow

DAGs

Data Profiling

Browse

Admin

Docs

About

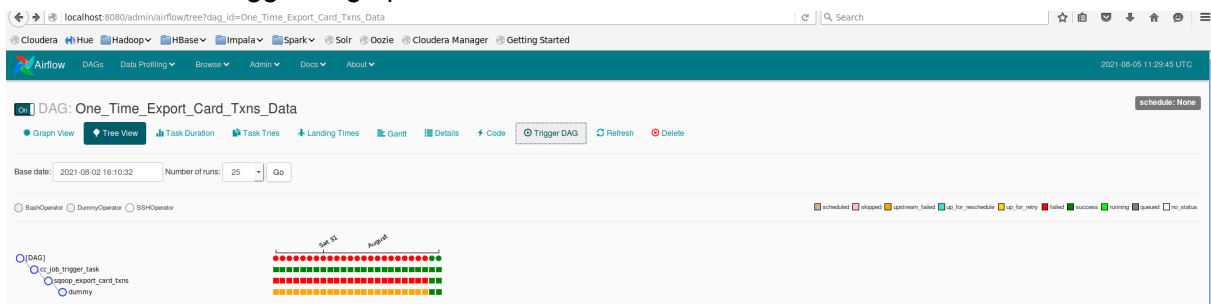
2021-08-05 11:26:18 UTC

DAGs

Search:


	DAG	Schedule	Owner	Recent Tasks	Last Run	DAG Runs	Links
	One_Time_Export_Cases_Texas_Data	None	airflow		2021-08-02 16:10		
	example_bash_operator	@daily	Airflow				
	example_branch_dag_operator_v3	@daily	Airflow				
	example_branch_operator	@daily	Airflow				

and click on the Trigger dag option



[←](#)
[localhost:8080/admin/airflow/trigger?dag_id=One_Time_Export_Card_Txns_Data&origin=%2Fadmin%2Ffairflow%2Ftree%3Fdag_id%](#)

[Cloudera](#)
[Hue](#)
[Hadoop](#)
[HBase](#)
[Impala](#)
[Spark](#)
[Solr](#)
[Oozie](#)
[Cloudera Manager](#)
[Getting Started](#)


[Airflow](#)
[DAGs](#)
[Data Profiling](#)
[Browse](#)
[Admin](#)
[Docs](#)
[About](#)

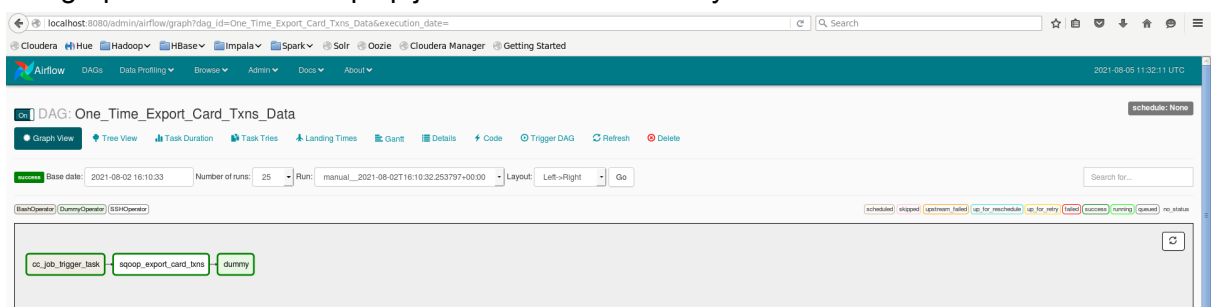
Trigger DAG: One_Time_Export_Card_Txns_Data

Configuration JSON (Optional)

To access configuration in your DAG use `{{ dag_run.conf }}`.

[Trigger](#)
[bail.](#)

The graph view after the sqoop job has run successfully



[ERROR SOLUTIONS:

1.If you run into errors while triggering the DAG like -> no such file / directory
Run the following to get rid of any special characters in the script

```
cat sqoop_export_card_txns.sh | tr -d '\r' > a.sh.ne
```

Then, rename a.sh.ne to sqoop_export_card_txns.sh back again:

```
mv a.sh.ne sqoop_export_card_txns.sh
```

2.If -> ERROR - SSH operator error: error running cmd: cd Desktop &&
./sqoop_import_card_txns.sh quickstart.cloudera:3306 bigdatapipeline root card_transactions,
error: 21/08/14 02:27:21 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6-cdh5.13.0:


It is due to java version, so install java 7 from link:

<https://www.oracle.com/ae/java/technologies/javase/javase7-archive-downloads.html>

and edit the bashrc file. Refer below ss:

-Execute below command in terminal to edit .bashrc file

```
gedit .bashrc
```



```
# .bashrc

# User specific aliases and functions

alias rm='rm -i'
alias cp='cp -i'
alias mv='mv -i'

export JAVA_HOME=/home/cloudera/Downloads/jdk1.7.0_76
export PATH=$JAVA_HOME/bin:$PATH

# Source global definitions
if [ -f /etc/bashrc ]; then
    . /etc/bashrc
fi
```

-Execute below command to bring the changes to effect:

```
source .bashrc
```

And now check java version, it should be java 1.7

]

--Verify count (mysql terminal)

```
select count(*) from stg_card_transactions;
```


--Remove Dups from Stg Table (mysql terminal)

Alter ignore table stg_card_transactions

Add unique index idx_card_txns(card_id,transaction_dt);

--Verify no dups

select card_id,transaction_dt,count(*) from stg_card_transactions group by
card_id,transaction_dt having count(*) >1;

--Dropping index used for removing dups

alter table stg_card_transactions drop index idx_card_txns;

--Loading main table

insert into card_transactions

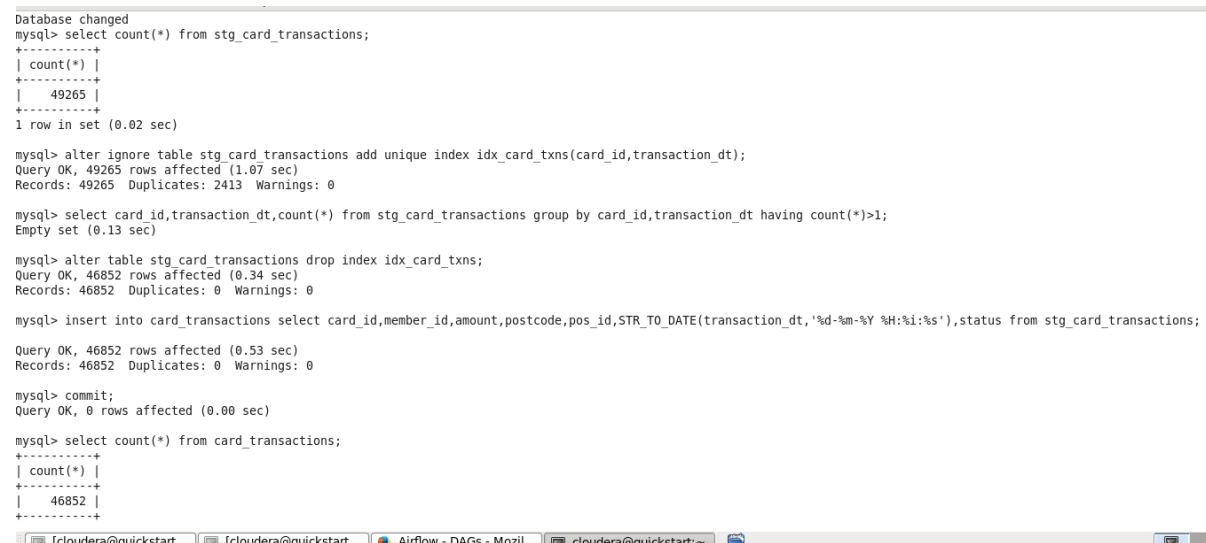
select

card_id,member_id,amount,postcode,pos_id,STR_TO_DATE(transaction_dt,'%d-%m-%Y
%H:%i:%s'),status from stg_card_transactions;

commit;

--Verify the count

select count(*) from card_transactions;



```
Database changed
mysql> select count(*) from stg_card_transactions;
+-----+
| count(*) |
+-----+
| 49265 |
+-----+
1 row in set (0.02 sec)

mysql> alter ignore table stg_card_transactions add unique index idx_card_txns(card_id,transaction_dt);
Query OK, 49265 rows affected (1.07 sec)
Records: 49265 Duplicates: 2413 Warnings: 0

mysql> select card_id,transaction_dt,count(*) from stg_card_transactions group by card_id,transaction_dt having count(*)>1;
Empty set (0.13 sec)

mysql> alter table stg_card_transactions drop index idx_card_txns;
Query OK, 46852 rows affected (0.34 sec)
Records: 46852 Duplicates: 0 Warnings: 0

mysql> insert into card_transactions select card_id,member_id,amount,postcode,pos_id,STR_TO_DATE(transaction_dt,'%d-%m-%Y %H:%i:%s'),status from stg_card_transactions;
Query OK, 46852 rows affected (0.53 sec)
Records: 46852 Duplicates: 0 Warnings: 0

mysql> commit;
Query OK, 0 rows affected (0.00 sec)

mysql> select count(*) from card_transactions;
+-----+
| count(*) |
+-----+
| 46852 |
+-----+
```

--Deleting the card_transactions file from HDFS

hadoop fs -rm /user/cloudera/project_input_data/card_transactions.csv

--Create member_score and member_details directory in HDFS

hadoop fs -mkdir project_input_data/member_score

hadoop fs -mkdir project_input_data/member_details

Task4: Creating HIVE Tables (Run in hive terminal)

Connect to Hive using -> \$hive

And enable bucketing using the following in the hive terminal-
SET HIVE.ENFORCE.BUCKETING=TRUE;

```
create external table if not exists member_score
(
  member_id string,
  score float
)
row format delimited fields terminated by ','
stored as textfile
location '/project_input_data/member_score/';
```

```
create external table if not exists member_details
(
  card_id bigint,
  member_id bigint,
  member_joining_dt timestamp ,
  card_purchase_dt timestamp ,
  country string,
  city string,
  score float
)
row format delimited fields terminated by ','
stored as textfile
location '/project_input_data/member_details/';
```

--Member score bucketed table(8 buckets)

```
create table if not exists member_score_bucketed
(
  member_id string,
  score float
)
CLUSTERED BY (member_id) into 8 buckets;
```

--Member details bucketed table(8 buckets)

```
create table if not exists member_details_bucketed
(
  card_id bigint,
  member_id bigint,
  member_joining_dt timestamp ,
  card_purchase_dt timestamp ,
  country string,
  city string,
  score float
)
CLUSTERED BY (card_id) into 8 buckets;
```

Hive-Hbase card_transactions table creation : External & Bucketed tables

```
create external table if not exists card_transactions (
  card_id bigint,
  member_id bigint,
  amount float,
  postcode int,
  pos_id bigint,
  transaction_dt timestamp,
  status string
)
row format delimited fields terminated by ','
stored as textfile
location '/project_input_data/card_transactions/';
```

```
create table card_transactions_bucketed
(
  cardid_txnts string,
  card_id bigint,
  member_id bigint,
  amount float,
  postcode int,
  pos_id bigint,
  transaction_dt timestamp,
  status string
)
CLUSTERED by (card_id) into 8 buckets
STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
WITH
SERDEPROPERTIES("hbase.columns.mapping"=":key,trans_data:card_id,trans_data:member_id,trans_data:amount,trans_data:postcode,trans_data:pos_id,trans_data:transaction_dt,trans_data:Status") TBLPROPERTIES ("hbase.table.name" = "card_transactions");
```

Hive-Hbase card_lookup table creation : Bucketed tables

```
create table card_lookup
(
  member_id bigint,
  card_id bigint ,
  ucl float ,
  score float,
  last_txn_time timestamp,
  last_txn_zip string
)
CLUSTERED by (card_id) into 8 buckets
STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
WITH
SERDEPROPERTIES("hbase.columns.mapping"=":key,lkp_data:member_id,lkp_data:ucl,lkp_data:score, lkp_data:last_txn_time,lkp_data:last_txn_zip")
TBLPROPERTIES ("hbase.table.name" = "card_lookup");
```

Task5: Inserting Data into optimised hive tables

Run member_score & member_details Airflow Job to load data from AWS to hive tables

Script for connecting to AWS RDS and load data to hive member score external table file path

[member_score.py present in the Big Data Project Demo -> airflow scripts folder in the google drive link]

--Inserting into member_score_bucketed table(Run from hive terminal)

insert into table member_score_bucketed;

select * from member_score;

Script for connecting to AWS RDS and load data to hive member details external table file path

[member_details.py present in the Big Data Project Demo -> airflow scripts folder in the google drive link]

--Inserting into member_details_bucketed table(Run from hive terminal)

insert into table member_details_bucketed select * from member_details;

OR

--Manually add memberscore.csv and memberdetails.csv to HDFS using the put command as shown below.

[cardmember.csv and memberscore.csv file will be provided in the google drive link mentioned at the beginning of this document under Project Dataset folder]

```
hadoop fs -put Desktop/memberscore.csv project_input_data/member_score/
```

```
hadoop fs -put Desktop/cardmembers.csv project_input_data/member_details/
```

--Load the data from hdfs to hive table (In Hive terminal)

```
load data inpath 'project_input_data/member_score/memberscore.csv' overwrite into table member_score;
```

```
load data inpath 'project_input_data/member_details/cardmembers.csv' overwrite into table member_details;
```

--To check if the data is loaded

```
select count(*) from member_score;
```

```
select count(*) from member_details;
```

--Inserting data into bucketed tables

```
insert into table member_score_bucketed  
select * from member_score;
```

```
insert into table member_details_bucketed  
select * from member_details;
```

Airflow script for importing card transactions data to HDFS card_transactions external table file path from MySQL

-Run the import_card_txns.py (copy sqoop_import_card_txns.sh file to Desktop and give permission to this file: `chmod 777 sqoop_import_card_txns.sh`)

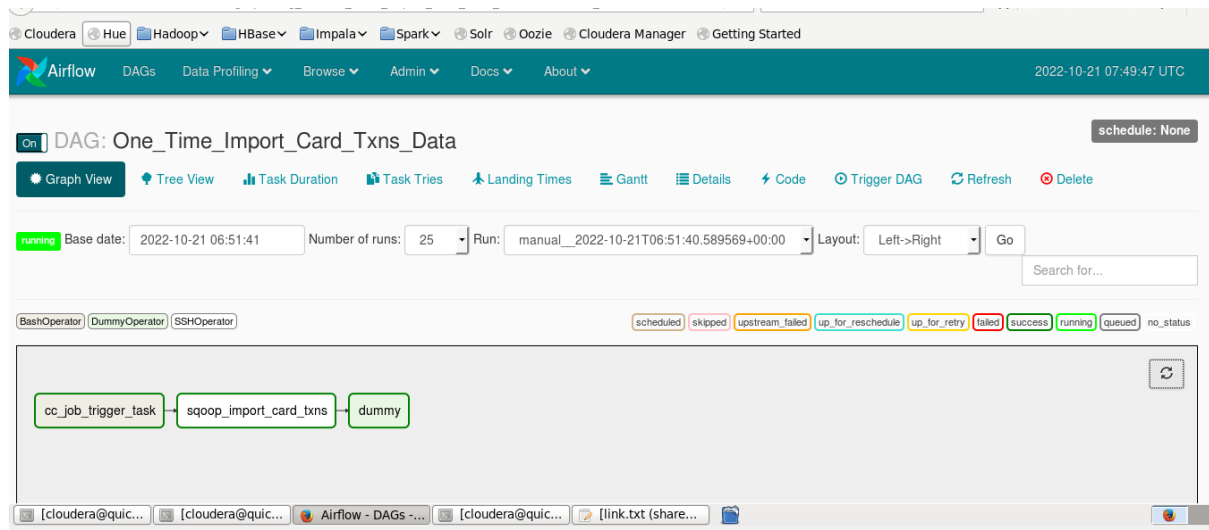
-Setting up Variable for import

click on **Variables**:

Under **create** tab, fill in the following details and save the connection

Key : card_txns_import_shell_command

Val : `./sqoop_import_card_txns.sh quickstart.cloudera:3306 bigdataproject root card_transactions`



ERROR SOLUTIONS:

if -> ERROR - SSH operator error: error running cmd: cd Desktop &&
 ./sqoop_import_card_txns.sh quickstart.cloudera:3306 bigdataprotect root card_transactions,
 error: 21/08/14 02:27:21 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6-cdh5.13.0
 Download java 7 and edit .bashrc file

```
.bashrc X
# .bashrc

# User specific aliases and functions

alias rm='rm -i'
alias cp='cp -i'
alias mv='mv -i'

export JAVA_HOME=/home/cloudera/Downloads/jdk1.7.0_76
export PATH=$JAVA_HOME/bin:$PATH

# Source global definitions
if [ -f /etc/bashrc ]; then
    . /etc/bashrc
fi
```

and execute below command in terminal:
 source .bashrc

And now check java version, it should be java 1.7
]

--Load card_txns_bucketed table with concatenated row key (In the Hive terminal)

insert into table card_transactions_bucketed

```
select concat_ws('~',cast(card_id as string),cast(transaction_dt as string)) as
cardid_txnts,card_id,member_id,amount,postcode,pos_id,transaction_dt,status
from card_transactions;
```

select count(*) from card_transactions_bucketed;

```
hive> insert into table card_transactions_bucketed select concat_ws('~',cast(card_id as string),cast(transaction_dt as string)) as cardid_txnts,card_id,member_id,amount
,postcode,pos_id,transaction_dt,status from card_transactions;
Query ID = cloudera_20221021002323_c6f1eff3-aadc-442f-8990-c6ee3f09e8d9
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1666259203287_0013, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1666259203287_0013/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1666259203287_0013
Hadoop job information for Stage-0: number of mappers: 2; number of reducers: 0
2022-10-21 00:23:39,812 Stage-0 map = 0%, reduce = 0%
2022-10-21 00:24:00,283 Stage-0 map = 50%, reduce = 0%, Cumulative CPU 2.66 sec
2022-10-21 00:24:11,493 Stage-0 map = 64%, reduce = 0%, Cumulative CPU 10.14 sec
2022-10-21 00:24:15,880 Stage-0 map = 100%, reduce = 0%, Cumulative CPU 11.23 sec
MapReduce Total cumulative CPU time: 11 seconds 230 msec
Ended Job = job_1666259203287_0013
MapReduce Jobs Launched:
Stage-Stage-0: Map: 2 Cumulative CPU: 11.23 sec HDFS Read: 4354793 HDFS Write: 0 SUCCESS
Total MapReduce CPU Time Spent: 11 seconds 230 msec
OK
Time taken: 56.935 seconds
hive> select count(*) from card_transactions_bucketed;
Query ID = cloudera_20221021002626_50068246-82d7-4abb-b286-c874e94b0461
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1666259203287_0014, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1666259203287_0014/
hives> select count(*) from card_transactions_bucketed;
Query ID = cloudera_20221021002626_50068246-82d7-4abb-b286-c874e94b0461
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1666259203287_0014, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1666259203287_0014/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1666259203287_0014
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2022-10-21 00:26:57,727 Stage-1 map = 0%, reduce = 0%
2022-10-21 00:27:15,916 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.24 sec
2022-10-21 00:27:26,911 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 5.87 sec
MapReduce Total cumulative CPU time: 5 seconds 870 msec
Ended Job = job_1666259203287_0014
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 5.87 sec HDFS Read: 16523 HDFS Write: 6 SUCCESS
Total MapReduce CPU Time Spent: 5 seconds 870 msec
OK
46852
```

ERROR SOLUTIONS:

Hbase connection refused error ->

Check if the Hbase server failed to start by

sudo service --status-all

To restart the Hbase server

sudo service hbase-master start

sudo service hbase-regionserver start]

Task6: Batch Processing

--Hbase Search functionality based on rowkey

scan 'card_transactions', {FILTER => "(PrefixFilter('340028465709212'))"}

```
hbase(main):001:0> scan 'card_transactions',{FILTER=>"(PrefixFilter('340028465709212'))"}
ROW                                COLUMN+CELL
340028465709212-2016-02-01 19:19:41 column=trans_data:Status, timestamp=1666337041817, value=GENUINE
340028465709212-2016-02-01 19:19:41 column=trans_data:amount, timestamp=1666337041817, value=301091.0
340028465709212-2016-02-01 19:19:41 column=trans_data:card_id, timestamp=1666337041817, value=340028465709212
340028465709212-2016-02-01 19:19:41 column=trans_data:member_id, timestamp=1666337041817, value=9250698176266
340028465709212-2016-02-01 19:19:41 column=trans_data:pos_id, timestamp=1666337041817, value=680905183957291
340028465709212-2016-02-01 19:19:41 column=trans_data:postcode, timestamp=1666337041817, value=15650
340028465709212-2016-02-01 19:19:41 column=trans_data:transaction_dt, timestamp=1666337041817, value=2016-02-01 19:19:41
340028465709212-2016-02-08 02:18:34 column=trans_data:Status, timestamp=1666337041817, value=GENUINE
340028465709212-2016-02-08 02:18:34 column=trans_data:amount, timestamp=1666337041817, value=451162.0
340028465709212-2016-02-08 02:18:34 column=trans_data:card_id, timestamp=1666337041817, value=340028465709212
340028465709212-2016-02-08 02:18:34 column=trans_data:member_id, timestamp=1666337041817, value=9250698176266
340028465709212-2016-02-08 02:18:34 column=trans_data:pos_id, timestamp=1666337041817, value=633734711060942
340028465709212-2016-02-08 02:18:34 column=trans_data:postcode, timestamp=1666337041817, value=71363
340028465709212-2016-02-08 02:18:34 column=trans_data:transaction_dt, timestamp=1666337041817, value=2016-02-08 02:18:34
340028465709212-2016-02-29 19:43:20 column=trans_data:Status, timestamp=1666337041817, value=GENUINE
340028465709212-2016-02-29 19:43:20 column=trans_data:amount, timestamp=1666337041817, value=107640.0
340028465709212-2016-02-29 19:43:20 column=trans_data:card_id, timestamp=1666337041817, value=340028465709212
340028465709212-2016-02-29 19:43:20 column=trans_data:member_id, timestamp=1666337041817, value=9250698176266
340028465709212-2016-02-29 19:43:20 column=trans_data:pos_id, timestamp=1666337041817, value=377613506830624
340028465709212-2016-02-29 19:43:20 column=trans_data:postcode, timestamp=1666337041817, value=43812
340028465709212-2016-02-29 19:43:20 column=trans_data:transaction_dt, timestamp=1666337041817, value=2016-02-29 19:43:20
340028465709212-2016-03-01 15:02:02 column=trans_data:Status, timestamp=1666337041817, value=GENUINE
340028465709212-2016-03-01 15:02:02 column=trans_data:amount, timestamp=1666337041817, value=983262.0
340028465709212-2016-03-01 15:02:02 column=trans_data:card_id, timestamp=1666337041817, value=340028465709212
340028465709212-2016-03-01 15:02:02 column=trans_data:member_id, timestamp=1666337041817, value=9250698176266
340028465709212-2016-03-01 15:02:02 column=trans_data:pos_id, timestamp=1666337041817, value=325960823570918
340028465709212-2016-03-01 15:02:02 column=trans_data:postcode, timestamp=1666337041817, value=13485
340028465709212-2016-03-01 15:02:02 column=trans_data:transaction_dt, timestamp=1666337041817, value=2016-03-01 15:02:02
340028465709212-2016-04-08 12:26:02 column=trans_data:Status, timestamp=1666337041817, value=GENUINE
340028465709212-2016-04-08 12:26:02 column=trans_data:amount, timestamp=1666337041817, value=3433123.0
340028465709212-2016-04-08 12:26:02 column=trans_data:card_id, timestamp=1666337041817, value=340028465709212
```

-Setting up Scala IDE

In your .bashrc file have the below 2 lines

```
=====
export JAVA_HOME=/home/cloudera/Downloads/jdk1.8.0_271
export PATH=$JAVA_HOME/bin:$PATH
```

Check java version using

```
=====
java -version
it should show java 1.8
```

Download the scala ide

```
=====
http://downloads.typesafe.com/scalaide-pack/4.7.0-vfinal-oxygen-212-20170929/scala-SDK-4.7.0-vfinal-2.12-linux.gtk.x86\_64.tar.gz and then extract this.
```

In your eclipse.ini file put the below line at the end.

```
=====
```

```
-Dosgi.bundles=org.eclipse.equinox.simpleconfigurator@1:start,org.eclipse.equinox.common@2:start,org.eclipse.equinox.ds@2:start,org.eclipse.equinox.event@2:start,org.eclipse.update.configurator@3:start,org.eclipse.core.runtime@start
```

Steps to start Scala IDE

```
=====
cd /home/cloudera/Downloads/eclipse
```


./eclipse -clean

Download spark 3.2.2 binaries in your cloudera using below link:

<https://archive.apache.org/dist/spark/>

Download Hbase Dependencies in your cloudera using below link:

http://codeload.github.com/gvreddy1210/Hbase_jars/zip/master

[if you download spark 3.2.2

you need to use scala 2.12 compiler]

-Create a new Scala Project (name:bigdataproject)

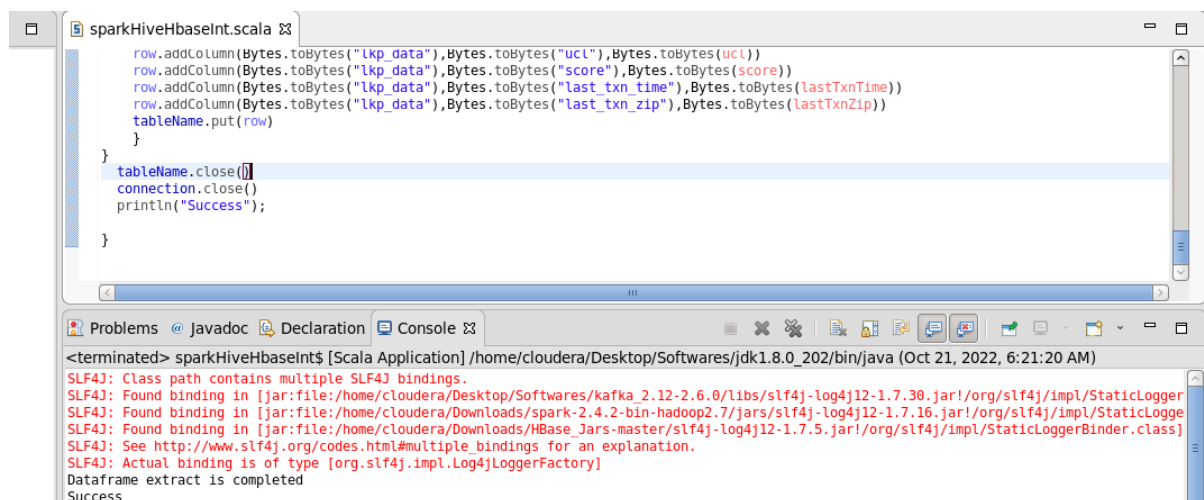
-Create new Scala Object(name:sparkHiveHbaseInt.scala) [Download the file

“sparkHiveHbaseInt.scala” from Code and scripts>Batch

processing>IntegrationCodes>sparkHiveHbaseInt]

-Add the required spark dependencies - “spark-3.2.2-bin-hadoop2.7” and “Hbase master”. (link is given above)

-Run as scala application sparkHiveHbaseInt.scala optimized scala code from Eclipse



```
sparkHiveHbaseInt.scala
row.addColumn(Bytes.toBytes("lqp_data"), Bytes.toBytes("ucl"), Bytes.toBytes(ucl))
row.addColumn(Bytes.toBytes("lqp_data"), Bytes.toBytes("score"), Bytes.toBytes(score))
row.addColumn(Bytes.toBytes("lqp_data"), Bytes.toBytes("last_txn_time"), Bytes.toBytes(lastTxnTime))
row.addColumn(Bytes.toBytes("lqp_data"), Bytes.toBytes("last_txn_zip"), Bytes.toBytes(lastTxnZip))
tableName.put(row)
}
tableName.close()
connection.close()
println("Success");
}
```

```
<terminated> sparkHiveHbaseInt$ [Scala Application] /home/cloudera/Desktop/Softwares/jdk1.8.0_202/bin/java (Oct 21, 2022, 6:21:20 AM)
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/cloudera/Desktop/Softwares/kafka_2.12-2.6.0/libs/slf4j-log4j12-1.7.30.jar!/org/slf4j/impl/StaticLogger
SLF4J: Found binding in [jar:file:/home/cloudera/Downloads/spark-2.4.2-bin-hadoop2.7/jars/slf4j-log4j12-1.7.16.jar!/org/slf4j/impl/StaticLogge
SLF4J: Found binding in [jar:file:/home/cloudera/Downloads/HBase_Jars-master/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
Dataframe extract is completed
Success
```

--Verify lookup table data from HBase

scan 'card_lookup'

```

hbase(main):001:0> scan 'card_lookup'
ROW                                COLUMN+CELL
340028465709212                    column=lkp_data:last_txn_time, timestamp=1666358522602, value=2018-01-02 03:25:35.0
340028465709212                    column=lkp_data:last_txn_zip, timestamp=1666358522602, value=24658
340028465709212                    column=lkp_data:member_id, timestamp=1666358522602, value=9250698176266
340028465709212                    column=lkp_data:score, timestamp=1666358522602, value=233.0
340028465709212                    column=lkp_data:ucl, timestamp=1666358522602, value=3.34787624E8
340054675199675                    column=lkp_data:last_txn_time, timestamp=1666358521295, value=2018-01-15 19:43:23.0
340054675199675                    column=lkp_data:last_txn_zip, timestamp=1666358521295, value=50140
340054675199675                    column=lkp_data:member_id, timestamp=1666358521295, value=835873341185231
340054675199675                    column=lkp_data:score, timestamp=1666358521295, value=631.0
340054675199675                    column=lkp_data:ucl, timestamp=1666358521295, value=1.4652636E7
340082915339645                    column=lkp_data:last_txn_time, timestamp=1666358513825, value=2018-01-26 19:03:47.0
340082915339645                    column=lkp_data:last_txn_zip, timestamp=1666358513825, value=17844
340082915339645                    column=lkp_data:member_id, timestamp=1666358513825, value=512969555857346
340082915339645                    column=lkp_data:score, timestamp=1666358513825, value=407.0
340082915339645                    column=lkp_data:ucl, timestamp=1666358513825, value=1.5726072E7
340134186926007                    column=lkp_data:last_txn_time, timestamp=1666358513620, value=2018-01-18 23:12:50.0
340134186926007                    column=lkp_data:last_txn_zip, timestamp=1666358513620, value=67576
340134186926007                    column=lkp_data:member_id, timestamp=1666358513620, value=887711945571282
340134186926007                    column=lkp_data:score, timestamp=1666358513620, value=9.5326131E8
340134186926007                    column=lkp_data:ucl, timestamp=1666358513620, value=1.6647037E7

```

--Verify lookup table data from Hive
select * from card_lookup limit 5;

```

hive> select * from card_lookup limit 5;
OK
340028465709212 9250698176266 3.34787616E8 233.0 2018-01-02 03:25:35 24658
340054675199675 835873341185231 1.4652636E7 631.0 2018-01-15 19:43:23 50140
340082915339645 512969555857346 1.5726072E7 407.0 2018-01-26 19:03:47 17844
340134186926007 887711945571282 9.5326131E8 614.0 2018-01-18 23:12:50 67576
340265728490548 680324265406190 1.6647037E7 202.0 2018-01-21 02:07:35 72435
Time taken: 0.323 seconds, Fetched: 5 row(s)
hive> █

```

Or

You can also trigger the same job through AIRFLOW as well

Airflow script to RUN the above JAR, batch_job

“cc_batch_job.py”

(Downloadable from drive) -> copy this file to the DAGS folder of Airflow and trigger the DAG once it is listed in the DAGS list

[Errors & Solutions

If error is - The root scratch dir: /tmp/hive on HDFS should be writable.

```

at scala.collection.immutable.List.foreach(List.scala:389)
at scala.App.main(App.scala:76)
at scala.App.main$(App.scala:74)
at sparkHiveHbaseInt$.main(sparkHiveHbaseInt.scala:20)
at sparkHiveHbaseInt$.main(sparkHiveHbaseInt.scala)
Caused by: java.lang.RuntimeException: The root scratch dir: /tmp/hive on HDFS should be writable. Current permissions are: rwx-----
at org.apache.hadoop.hive.ql.session.SessionState.start(SessionState.java:522)
at org.apache.spark.sql.hive.client.HiveClientImpl.newState(HiveClientImpl.scala:193)
at org.apache.spark.sql.hive.client.HiveClientImpl.<init>(HiveClientImpl.scala:117)
at sun.reflect.NativeConstructorAccessorImpl.newInstance0(Native Method)
at sun.reflect.NativeConstructorAccessorImpl.newInstance(NativeConstructorAccessorImpl.java:62)
at sun.reflect.DelegatingConstructorAccessorImpl.newInstance(DelegatingConstructorAccessorImpl.java:45)
at java.lang.reflect.Constructor.newInstance(Constructor.java:423)
at org.apache.spark.sql.hive.client.IsolatedClientLoader.createClient(IsolatedClientLoader.scala:271)
at org.apache.spark.sql.hive.HiveUtils$.newClientForMetadata(HiveUtils.scala:384)
at org.apache.spark.sql.hive.HiveUtils$.newClientForMetadata(HiveUtils.scala:286)
at org.apache.spark.sql.hive.HiveExternalCatalog.clientsToCompute(HiveExternalCatalog.scala:66)
at org.apache.spark.sql.hive.HiveExternalCatalog.client(HiveExternalCatalog.scala:65)
at org.apache.spark.sql.hive.HiveExternalCatalog.$anonfun$databaseExists$(HiveExternalCatalog.scala:215)
at scala.runtime.java8.JFunction0$mcZ$sp.apply(JFunction0$mcZ$sp.java:12)
at org.apache.spark.sql.hive.HiveExternalCatalog.withClient(HiveExternalCatalog.scala:97)
... 93 more
Caused by: java.lang.RuntimeException: The root scratch dir: /tmp/hive on HDFS should be writable. Current permissions are: rwx-----
at org.apache.hadoop.hive.ql.session.SessionState.createRootHDFSDir(SessionState.java:612)
at org.apache.hadoop.hive.ql.session.SessionState.createSessionDirs(SessionState.java:554)
at org.apache.hadoop.hive.ql.session.SessionState.start(SessionState.java:508)
... 107 more

```

Solution - Rebuild the project by adding the spark binaries - spark-3.2.2-bin-hadoop2.7]

Task7: Stream Processing

- Download the file "classReadFromKafka.scala" (from Code and Scripts>Streaming>project_repo_streaming>streamjobClasses)
- Download the file "classCardValidation.scala" (from Code and Scripts>Streaming>project_repo_streaming>streamjobClasses)
- Create new scala project (name:bigdatastreaming)
- Add the above 2 scala files under src folder of the "bigdatastreaming"

Download spark 2.4.3 binaries in your cloudera using below link:

<https://archive.apache.org/dist/spark/>

Here, set the scala compiler: 2.11

-Add the below jars:

1. "kafka-spark(dependencies)" folder (from Code and Scripts>Streaming>project_repo_streaming>kafka-spark(dependencies))
2. Add the external jars from the below link for Hbase:
http://codeload.github.com/gvreddy1210/Hbase_jars/zip/master
3. Add distanceFindjar(from Code and

Scripts>Streaming>project_repo_streaming>distanceFindjar)

[Note: Download the zipCodePosId.csv from downloadables and keep it in your Desktop, and change the path of this csv file in the code as per your path in class CardValidation.scala code]

-Starting Kafka Producer and Consumer and creating the KafkaTopic

--Start Kafka

```
cd /home/cloudera/Desktop/Softwares/kafka_2.12-2.6.0/bin
```

```
./kafka-server-start.sh ../config/server.properties
```

--Create Topic(In a new terminal)

```
cd /home/cloudera/Desktop/Softwares/kafka_2.12-2.6.0/bin
```

```
./kafka-topics.sh --create --topic cctxnstopic --bootstrap-server localhost:9092 --partitions 1
```

[Note: Use the topic name created here itself in your code also]

--Kafka Producer

```
./kafka-console-producer.sh --broker-list localhost:9092 --topic cctxnstopic
```

--Sample records to pass

```
//genuine
{"card_id":555059812846420,"member_id":6460955484292953,"amount":9000,"pos_id":4444,"post_code":10537,"transc_dt":"2021-03-09 07:28:43"}
```

```
//fraud
{"card_id":555059812846420,"member_id":6460955484292953,"amount":9000000000,"pos_id":4444,"post_code":10537,"transc_dt":"2021-03-09 07:28:43"}
```

Note : the amount and transac_dt has to be changed to test for Fraud transactions.

-Run as scala application classReadFromKafka.scala optimised scala code from Eclipse

Output after producing genuine transaction:

```
Batchid : 1
From ValidationClass:[Ljava.lang.String;@6aa942a6
Tables connected
Row Fetched From Hbase
Values extracted to local from hbase
Bytes converted into String591.0
Validation 1 is done
txnAmountVal :9000.0
uclVal :9644359.0
Validation 2 is done
Distance Extracted :15.918489714172177
Time conversion done
time difference in hours :-1
velocity :-15.918489714172177
Validation is completed
memberIdVal:6460955484292953
uclVal:9644359.0
scoreVal:591.0
last_txn_timeVal:2021-03-09 09:00:00
postCodeVal:10537
Cardlookup is updated
CardTransactions is updated
=====GENUINE=====
success :1
```

Output after producing fraud transaction:

```

Batchid : 2
From ValidationClass:[Ljava.lang.String;@4540b9d8
Tables connected
Row Fetched From Hbase
Values extracted to local from hbase
Bytes converted into String591.0
Validation 1 is done
txnAmountVal :9.0E9
uclVal :9644359.0
Validation 2 is done
Distance Extracted :15.918489714172177
Time conversion done
time difference in hours :0
velocity :Infinity
Validation is completed
memberIdVal:6460955484292953
uclVal:9644359.0
scoreVal:591.0
last_txn_timeVal:2021-03-09 07:28:43
postCodeVal:10537
Cardlookup is updated
CardTransactions is updated
=====FRAUD=====
success :2

```

Verification in card_transaction table to check if its updated:

Run following command in Hbase shell:

```
hbase(main):007:0> scan 'card_transactions',{FILTER=>(PrefixFilter('555059812846420'))}
```

5550598128464202021-03-09 09:00:00	column=trans_data:amount, timestamp=1671617540782, value=9000.0
5550598128464202021-03-09 09:00:00	column=trans_data:card_id, timestamp=1671617540782, value=555059812846420
5550598128464202021-03-09 09:00:00	column=trans_data:member_id, timestamp=1671617540782, value=6460955484292953
5550598128464202021-03-09 09:00:00	column=trans_data:pos_id, timestamp=1671617540782, value=4444
5550598128464202021-03-09 09:00:00	column=trans_data:postcode, timestamp=1671617540782, value=10537
5550598128464202021-03-09 09:00:00	column=trans_data:status, timestamp=1671617540782, value=GENIUNE
5550598128464202021-03-09 09:00:00	column=trans_data:transaction_dt, timestamp=1671617540782, value=2021-03-09 09:00:00

5 row(s) in 0.1670 seconds

```
hbase(main):008:0> █
```

5550598128464202021-03-09 07:28:43	column=trans_data:amount, timestamp=1671617894228, value=9.0E9
5550598128464202021-03-09 07:28:43	column=trans_data:card_id, timestamp=1671617894228, value=555059812846420
5550598128464202021-03-09 07:28:43	column=trans_data:member_id, timestamp=1671617894228, value=6460955484292953
5550598128464202021-03-09 07:28:43	column=trans_data:pos_id, timestamp=1671617894228, value=4444
5550598128464202021-03-09 07:28:43	column=trans_data:postcode, timestamp=1671617894228, value=10537
5550598128464202021-03-09 07:28:43	column=trans_data:status, timestamp=1671617894228, value=FRAUD
5550598128464202021-03-09 07:28:43	column=trans_data:transaction_dt, timestamp=1671617894228, value=2021-03-09 07:28:43

card_transactions table is updated!!

=====

